


# TYPE ENFORCEMENT

## *PROCESS TYPES*

The SELinux primary model of enforcement is called type enforcement. Basically, this means we define the label on a process based on its type, and the label on a file system object based on its type.

Imagine a system where we define types on objects like cats and dogs. A cat and dog are process types.



CAT

DOG

## OBJECT TYPES

We have a class of objects that they want to interact with which we call food. And I want to add types to the food, cat\_chow and dog\_chow.



CAT\_CHOW




DOG\_CHOW

## POLICY RULES

As a policy writer, I would say that a dog has permission to eat dog\_chow food and a cat has permission to eat cat\_chow food. In SELinux we would write this rule in policy, as shown below:




ALLOW



CAT

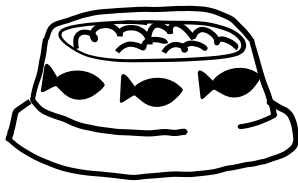


ALLOW



DOG





CAT\_CHOW:FOOD

+



EAT





DOG\_CHOW:FOOD

+




EAT

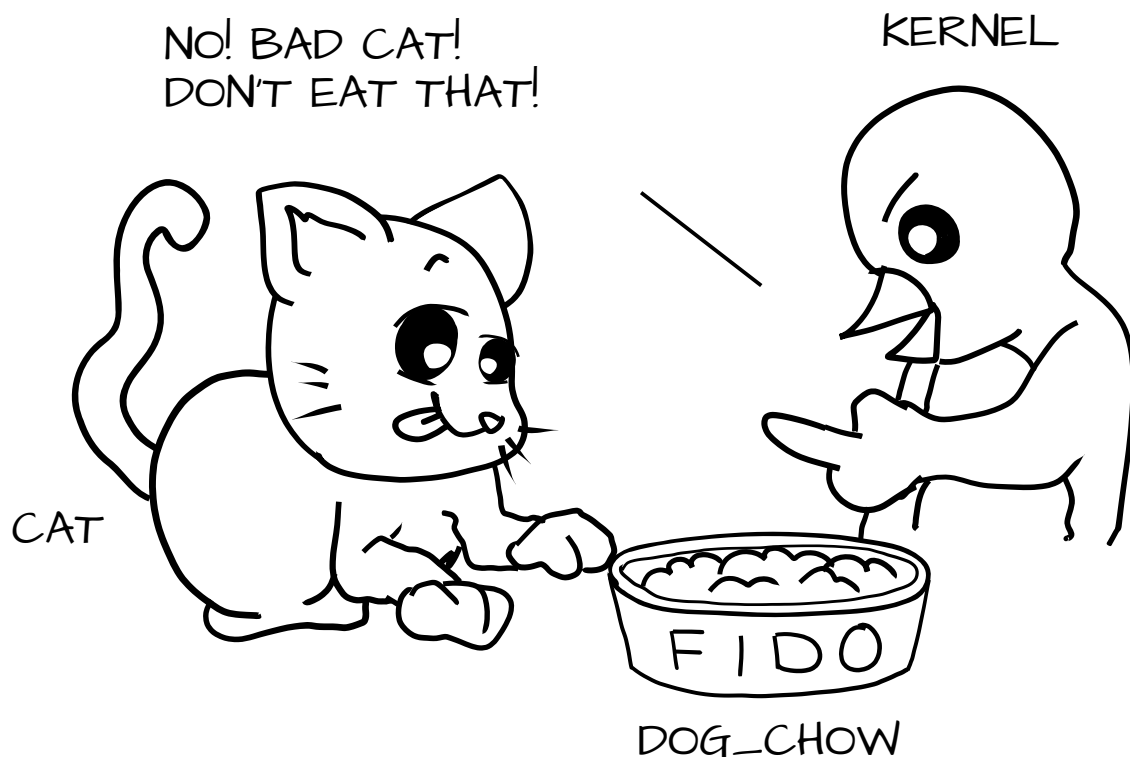
With these rules the kernel would allow the cat process to eat food labeled cat\_chow and the dog to eat food labeled dog\_chow.



But in an SELinux system, everything is denied by default. This means that if the dog process tried to eat the cat\_chow, the kernel would prevent it.

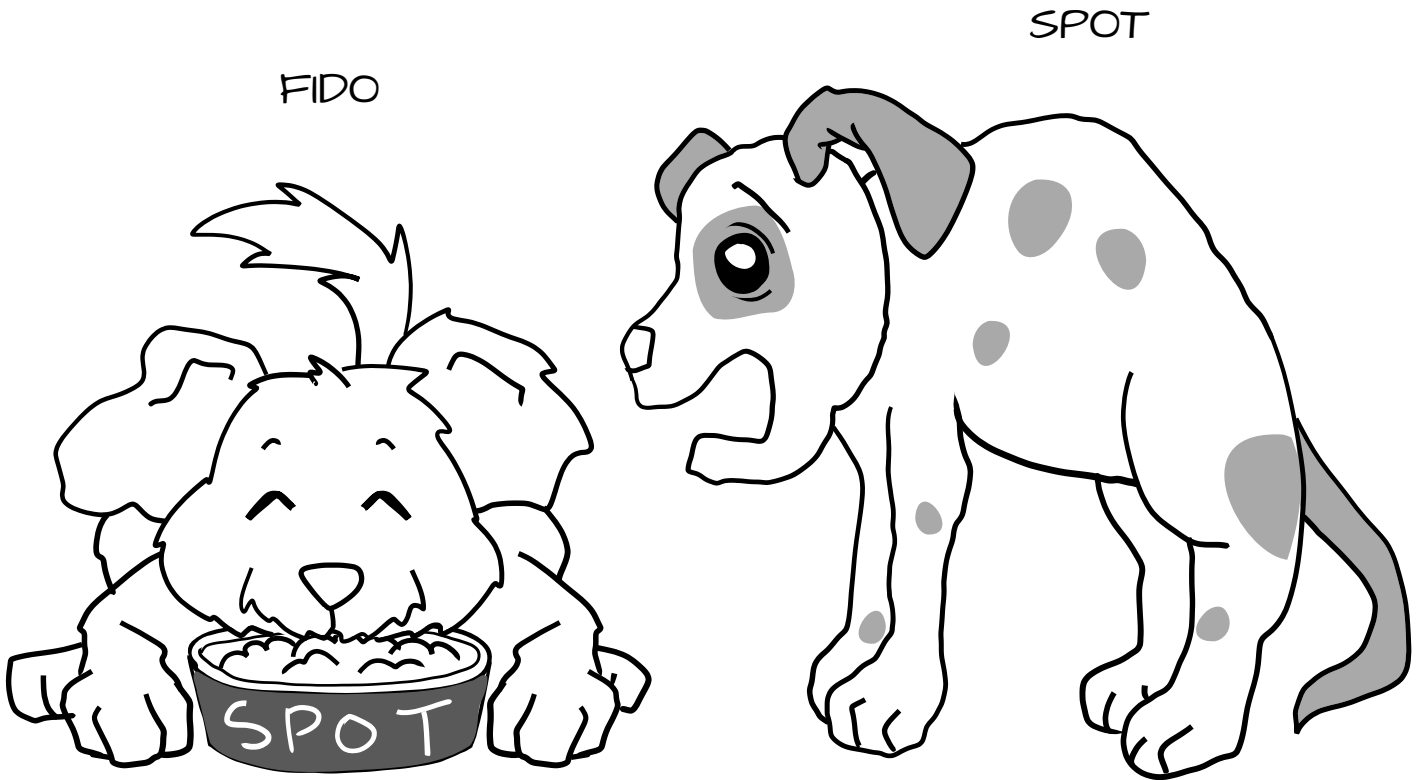


Likewise, cats would not be allowed to touch dog food.



# MCS ENFORCEMENT


We've typed the dog process and cat process, but what happens if you have multiple dog processes: Fido and Spot? You want to stop Fido from eating Spot's dog\_chow.



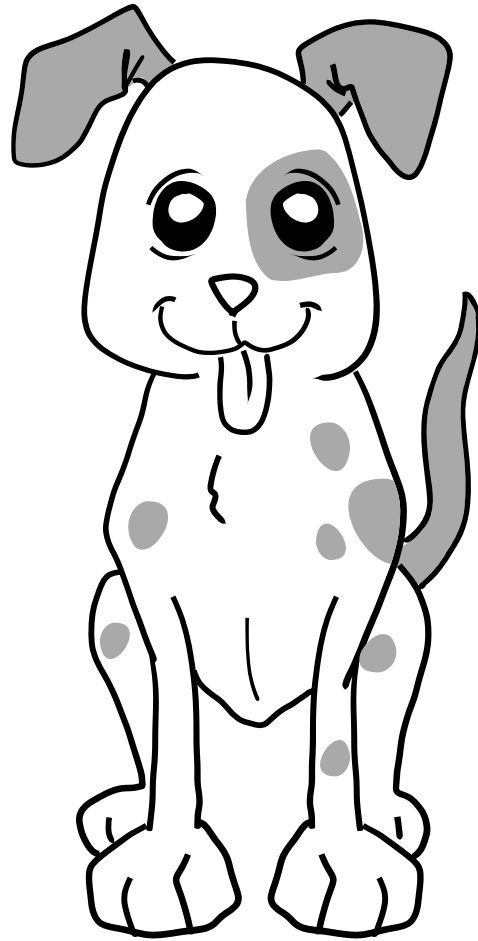
One solution would be to create lots of new types, like Fido\_dog and Fido\_dog\_chow. But, this will quickly become unruly because all dogs have pretty much the same permissions.

To handle this we developed a new form of enforcement, which we call Multi Category Security (MCS). In MCS, we add another section of the label which we can apply to the dog process and to the dog\_chow food. Now we label the dog process as dog:random1 (Fido) and dog:random2 (Spot).

We label the dog chow as dog\_chow:random1 (Fido) and dog\_chow:random2 (Spot).




DOG:RANDOM1



DOG:RANDOM2



DOG\_CHOW:  
RANDOM1




DOG\_CHOW:  
RANDOM2

MCS rules say that if the type enforcement rules are OK and the random MCS labels match exactly, then the access is allowed. If not, it is denied.

# TYPE ENFORCEMENT

Fido (dog:random1) trying to eat cat\_chow:food is denied by type enforcement.



---

Fido (dog:random1) is allowed to eat dog\_chow:random1.




DOG:RANDOM1

DOG\_CHOW:  
RANDOM1



# MCS ENFORCEMENT

Fido (dog:random1) denied to eat spot's (dog\_chow:random2) food.

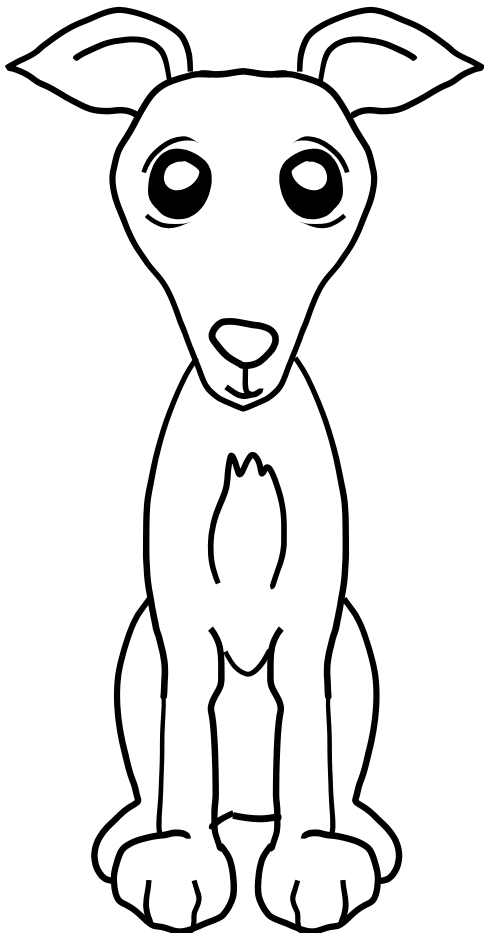


# MLS ENFORCEMENT

Another form of SELinux enforcement, used much less frequently, is called Multi Level Security (MLS); it was developed back in the 60s and is used mainly in trusted operating systems like Trusted Solaris.

The main idea is to control processes based on the level of the data they will be using: a secret process can not read top-secret data.

Instead of talking about different dogs, we now look at different breeds. We might have a Greyhound and a Chihuahua:




GREYHOUND




CHIHUAHUA

We might want to allow the Greyhound to eat any dog food, but a Chihuahua could choke if it tried to eat Greyhound dog food.

We want to label the Greyhound as dog:Greyhound and his dog food as dog\_chow:Greyhound, and label the Chihuahua as dog:Chihuahua and his food as dog\_chow:Chihuahua.




DOG\_CHOW:  
GREYHOUND



DOG\_CHOW:  
CHIHUAHUA

With the MLS policy, we would have the MLS Greyhound label dominate the Chihuahua label. This means dog:Greyhound is allowed to eat dog\_chow:Greyhound and dog\_chow:Chihuahua.




DOG:  
CHIHUAHUA



DOG\_CHOW:  
CHIHUAHUA

But dog:Chihuahua is not allowed to eat dog\_chow:Greyhound.



Of course, dog:Greyhound and dog:Chihuahua are still prevented from eating cat\_chow:Siamese by type enforcement, even if the MLS type Greyhound dominates Siamese.

